

Random Forest and Boosting

Cornell CS 4/5780 — Spring 2022

Random Forest

One of the most famous and useful bagged algorithms is the **Random Forest**! A Random Forest is essentially nothing else but **bagged** decision trees, with a slightly modified splitting criteria. The algorithm works as follows:

1. Sample m data sets D_1, \dots, D_m from D with replacement.
2. For each D_j train a full decision tree $h_j(\cdot)$ (max-depth= ∞) with one small modification: before each **split** **randomly subsample** $k \leq d$ features (without replacement) and only consider these for your split. (This further increases the variance of the trees.)
3. The final classifier is $h(\mathbf{x}) = \frac{1}{m} \sum_{j=1}^m h_j(\mathbf{x})$.

The Random Forest is one of the best, most popular and easiest to use out-of-the-box classifier. There are two reasons for this:

- The RF only has **two hyper-parameters**, m and k . It is extremely *insensitive* to both of these. A good choice for k is $k = \sqrt{d}$ (where d denotes the number of features). You can set m as large as you can afford.
- Decision trees do not require a lot of preprocessing. For example, the features can be of different scale, magnitude, or slope. This can be highly advantageous in scenarios with heterogeneous data, for example the medical settings where features could be things like *blood pressure, age, gender, ...*, each of which is recorded in completely different units.

Useful **variants** of Random Forests:

- Split each training set into two partitions $D_i = D_i^A \cup D_i^B$, where $D_i^A \cap D_i^B = \emptyset$. Build the tree on D_i^A and estimate the leaf labels on D_i^B . You must stop splitting if a leaf has only a single point in D_i^B in it. This has the advantage that each tree and also the RF classifier become consistent (this roughly means that it approaches the **Bayes optimal classifier** in the **limit of large training set**).
- Do not grow each tree to its full depth, instead **prune based on the leave out samples**. This can further improve your bias/variance trade-off.

Scenario: Hypothesis class \mathbb{H} , whose set of classifiers has large bias and the training error is high (e.g. CART trees with very limited depth.)

Famous question: In his machine learning class project in 1988 Michael Kearns **famously asked the question:** Can **weak learners** (H) be **combined** to generate a **strong learner** with **low bias**?

Famous answer: **Yes!** (Robert Schapire in 1990).

Solution: Create ensemble classifier $H_T(\vec{x}) = \sum_{t=1}^T \alpha_t h_t(\vec{x})$. This **ensemble** classifier is built in an **iterative** fashion. In iteration t we **add the classifier** $\alpha_t h_t(\vec{x})$ to the ensemble. At test time we evaluate all classifier and return the weighted sum.

The process of constructing such an ensemble in a stage-wise fashion is very similar to **gradient descent**. However, instead of updating the model parameters in each iteration, we add functions to our ensemble.

Let ℓ denote a (convex and differentiable) loss function. With a little abuse of notation we write

$$\ell(H) = \frac{1}{n} \sum_{i=1}^n \ell(H(\mathbf{x}_i), y_i).$$

Assume we have already finished t iterations and already have an ensemble classifier $H_t(\vec{x})$. Now in iteration $t + 1$ we want to **add one more** weak learner h_{t+1} to the ensemble. To this end we search for the weak learner that minimizes the loss the most,

$$h_{t+1} = \operatorname{argmin}_{h \in \mathbb{H}} \ell(H_t + \alpha h).$$

Once h_{t+1} has been found, we add it to our ensemble, i.e. $H_{t+1} := H_t + \alpha h$.

How can we find such $h \in \mathbb{H}$?

Answer: Use **gradient descent** in **function space**. In function space, inner product can be defined as $\langle h, g \rangle = \int_x h(x)g(x)dx$. Since we only have training set, we define $\langle h, g \rangle = \sum_{i=1}^n h(\mathbf{x}_i)g(\mathbf{x}_i)$.

Gradient descent in functional space

Given H , we want to find the step-size α and (weak learner) h to minimize the loss $\ell(H + \alpha h)$. Use **Taylor Approximation** on $\ell(H + \alpha h)$.

$$\ell(H + \alpha h) \approx \ell(H) + \alpha \langle \nabla \ell(H), h \rangle.$$

This approximation (of ℓ as a linear function) only holds within a small region around $\ell(H)$, i. as long as α is small. We therefore fix it to a small constant (e.g. $\alpha \approx 0.1$). With the step-size α fixed, we can use the approximation above to find an almost optimal h :

$$\operatorname{argmin}_{h \in \mathbb{H}} \ell(H + \alpha h) \approx \operatorname{argmin}_{h \in \mathbb{H}} \langle \nabla \ell(H), h \rangle = \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n \frac{\partial \ell}{\partial [H(\mathbf{x}_i)]} h(\mathbf{x}_i)$$

We can write $\ell(H) = \sum_{i=1}^n \ell(H(\mathbf{x}_i)) = \ell(H(x_1), \dots, H(x_n))$ (each prediction is an input to the loss function)

$$\frac{\partial \ell}{\partial H}(\mathbf{x}_i) = \frac{\partial \ell}{\partial [H(\mathbf{x}_i)]}$$

So we can do boosting if we have an algorithm \mathbb{A} to solve

$$h_{t+1} = \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n \underbrace{\frac{\partial \ell}{\partial [H(\mathbf{x}_i)]}}_{r_i} h(x)$$

We need a function $\mathbb{A}(\{(\mathbf{x}_1, r_1), \dots, (\mathbf{x}_n, r_n)\}) = \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n r_i h(\mathbf{x}_i)$. In order to make progress this h does not have to be great. We still make progress as long as $\sum_{i=1}^n r_i h(\mathbf{x}_i) < 0$.

Generic boosting (a.k.a Anyboost)

Case study #1: Gradient Boosted Regression Tree(GBRT)

- Classification ($y_i \in \{+1, -1\}$) or (even multi-dimensional) regression ($y_i \in \mathcal{R}^k$)
- Weak learners, $h \in \mathbb{H}$, are regressors, $h(\mathbf{x}) \in \mathcal{R}, \forall \mathbf{x}$, typically **fixed-depth** (e.g. depth=4) **regression trees** (hence the name).
- Step size α is **fixed** to a **small** constant (hyper-parameter).
- Loss function: Any differentiable convex loss that decomposes over the samples $\mathcal{L}(H) = \sum_{i=1}^n \ell(H(\mathbf{x}_i))$

```

Input:  $\ell, \alpha, \{(\mathbf{x}_i, y_i)\}, \mathbb{A}$ 
 $H_0 = 0$ 
for  $t=0:T-1$  do
   $\forall i: r_i = \frac{\partial \ell(H_t(\mathbf{x}_1), y_1), \dots, (H_t(\mathbf{x}_n), y_n))}{\partial H(\mathbf{x}_i)}$ 
   $h_{t+1} = \mathbb{A}(\{(\mathbf{x}_1, r_1), \dots, (\mathbf{x}_n, r_n)\}) = \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n r_i h(\mathbf{x}_i)$ 
  if  $\sum_{i=1}^n r_i h_{t+1}(\mathbf{x}_i) < 0$  then
     $H_{t+1} = H_t + \alpha_{t+1} h_{t+1}$ 
  else
    return ( $H_t$ )
  end
end
return  $H_T$ 

```

Algorithm 1: AnyBoost in Pseudo-Code

In order to use regression trees for **gradient boosting**, we must be able to find a tree $h()$ that maximizes $h = \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n r_i h(\mathbf{x}_i)$ where $r_i = \frac{\partial \ell}{\partial H(\mathbf{x}_i)}$.

We will make **two assumptions**:

- First, we assume that $\sum_{i=1}^n h^2(\mathbf{x}_i) = \text{constant}$. This is simple to do (we normalize the predictions) and important because we could always decrease $\sum_{i=1}^n h(\mathbf{x}_i) r_i$ by **rescaling** h with a large constant. By fixing $\sum_{i=1}^n h^2(\mathbf{x}_i)$ to a constant we are essentially fixing the vector h to lie on a **circle**, and we are only concerned with its **direction** but **not its length**.
- CART trees are negation closed, i.e. $\forall h \in \mathbb{H} \Rightarrow \exists -h \in \mathbb{H}$. (This is generally true.)
- We can define the negative gradient as $t_i = -r_i$.

$$\begin{aligned}
 & \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n r_i h(\mathbf{x}_i) \quad (\text{This is the original AnyBoost formulation.}) \\
 &= \operatorname{argmin}_{h \in \mathbb{H}} -2 \sum_{i=1}^n t_i h(\mathbf{x}_i) \quad (\text{Swapping in } t_i \text{ for } -r_i \text{ and multiplying by 2, which is a constant.}) \\
 &= \operatorname{argmin}_{h \in \mathbb{H}} \underbrace{\sum_{i=1}^n t_i^2}_{\text{constant}} - 2 \sum_{i=1}^n t_i h(\mathbf{x}_i) + \underbrace{(h(\mathbf{x}_i))^2}_{\text{constant}} \quad (\text{Adding constant } \sum_i t_i^2 + h(\mathbf{x}_i)^2.) \\
 &= \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n (h(\mathbf{x}_i) - t_i)^2
 \end{aligned}$$

In other words, we can **use the good old Regression trees** and **feed in the value t_i as labels for each \mathbf{x}_i** . Each iteration we **build a new tree** for a different set of "labels" t_1, \dots, t_n .

If the loss function ℓ is the squared loss, i.e. $\ell(H) = \frac{1}{2} \sum_{i=1}^n (H(\mathbf{x}_i) - y_i)^2$, then it is easy to show that

$$t_i = -\frac{\partial \ell}{\partial H(\mathbf{x}_i)} = y_i - H(\mathbf{x}_i),$$

which is simply the **residual**, i.e. \mathbf{r} is the vector pointing from \mathbf{y} to \mathbf{H} . However, it is important that you can use any other differentiable and convex loss function ℓ , and the solution for your next weak learner $h()$ will always be the regression tree minimizing the squared loss.

Case Study #2: AdaBoost

- Setting: Classification ($y_i \in \{+1, -1\}$)
- Weak learners: $h \in \mathbb{H}$ are binary, $h(\mathbf{x}_i) \in \{-1, +1\}, \forall \mathbf{x}$
- Step-size: We perform line-search to obtain best step-size α .
- Loss function: Exponential loss $\ell(H) = \sum_{i=1}^n e^{-y_i H(\mathbf{x}_i)}$

Finding the best weak learner

First we compute the gradient $r_i = \frac{\partial \ell}{\partial H(\mathbf{x}_i)} = -y_i e^{-y_i H(\mathbf{x}_i)}$.

For notational convenience (and for reason that will become clear in a little bit), let us define $w_i = \frac{1}{Z} e^{-y_i H(\mathbf{x}_i)}$, where $Z = \sum_{i=1}^n e^{-y_i H(\mathbf{x}_i)}$ is a **normalizing factor** so that $\sum_{i=1}^n w_i = 1$. *Note that the normalizing constant Z is identical to the loss function. Each weight w_i therefore has a very nice interpretation. It is the **relative contribution** of the training point (\mathbf{x}_i, y_i) towards the overall loss.*

In order to find the best next weak learner, we need to solve the optimization problem: (in the following, we will make use of the fact that $h(\mathbf{x}_i) \in \{+1, -1\}$.)

GBRT in Pseudo Code

```

Input:  $\ell, \alpha, \{(\mathbf{x}_i, y_i)\}, \mathbb{A}$ 
 $H = 0$ 
for  $t=1:T$  do
   $\forall i: t_i = y_i - H(\mathbf{x}_i)$ 
   $h = \operatorname{argmin}_{h \in \mathbb{H}} (h(\mathbf{x}_i) - t_i)^2$ 
   $H \leftarrow H + \alpha h$ 
end
return  $H$ 

```

$$\begin{aligned}
h(\mathbf{x}_i) &= \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i=1}^n r_i h(\mathbf{x}_i) && \left(\text{substitute in: } r_i = e^{-H(\mathbf{x}_i)y_i} \right) \\
&= \operatorname{argmin}_{h \in \mathbb{H}} - \sum_{i=1}^n y_i e^{-H(\mathbf{x}_i)y_i} h(\mathbf{x}_i) && \left(\text{substitute in: } w_i = \frac{1}{Z} e^{-H(\mathbf{x}_i)y_i} \right) \\
&= \operatorname{argmin}_{h \in \mathbb{H}} - \sum_{i=1}^n w_i y_i h(\mathbf{x}_i) && \left(y_i h(\mathbf{x}_i) \in \{+1, -1\} \text{ with } h(\mathbf{x}_i)y_i = 1 \iff h(\mathbf{x}_i) = y_i \right) \\
&= \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i - \sum_{i:h(\mathbf{x}_i) = y_i} w_i && \left(\sum_{i:h(\mathbf{x}_i) = y_i} w_i = 1 - \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i \right) \\
&= \operatorname{argmin}_{h \in \mathbb{H}} \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i && \left(\text{This is the weighted classification error.} \right)
\end{aligned}$$

Let us denote this **weighted classification error** as $\epsilon = \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i$. So for AdaBoost, we only need a classifier that can take **training** data and a distribution over the training set (i.e. normalized weights w_i for all training samples) and which returns a classifier $h \in H$ that **reduces the weighted classification error** of these training samples. It doesn't have to do all that well, in order for the inner-product $\sum_i r_i h(\mathbf{x}_i)$ to be **negative**, it just needs **less than $\epsilon < 0.5$ weighted training error**.

Finding the stepsize α

In the previous example, GBRT, we set the stepsize α to be a **small constant**. As it turns out, in the **AdaBoost** setting we can find the **optimal stepsize** (i.e. the one that minimizes ℓ the most) in **closed form** every time we take a "gradient" step.

When we are given ℓ, H, h , we would like to solve the following optimization problem:

$$\begin{aligned}
\alpha &= \operatorname{argmin}_{\alpha} \ell(H + \alpha h) \\
&= \operatorname{argmin}_{\alpha} \sum_{i=1}^n e^{-y_i [H(\mathbf{x}_i) + \alpha h(\mathbf{x}_i)]}
\end{aligned}$$

We differentiate w.r.t. α and equate with zero:

$$\begin{aligned}
\sum_{i=1}^n y_i h(\mathbf{x}_i) e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} &= 0 && \left(y_i h(\mathbf{x}_i) \in \{+1, -1\} \right) \\
- \sum_{i:h(\mathbf{x}_i)y_i=1} e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} + \sum_{i:h(\mathbf{x}_i)y_i=-1} e^{-(y_i H(\mathbf{x}_i) + \alpha y_i h(\mathbf{x}_i))} &= 0 && \left(w_i = \frac{1}{Z} e^{-y_i H(\mathbf{x}_i)} \right) \\
- \sum_{i:h(\mathbf{x}_i)y_i=1} w_i e^{-\alpha} + \sum_{i:h(\mathbf{x}_i)y_i=-1} w_i e^{\alpha} &= 0 && \left(\epsilon = \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i \right) \\
-(1 - \epsilon) e^{-\alpha} + \epsilon e^{\alpha} &= 0 \\
e^{2\alpha} &= \frac{1 - \epsilon}{\epsilon} \\
\alpha &= \frac{1}{2} \ln \frac{1 - \epsilon}{\epsilon}
\end{aligned}$$

It is unusual that we can find the optimal step-size in such a simple closed form. One consequence is that AdaBoost **converges extremely fast**.

Re-normalization

After you take a step, i.e. $H_{t+1} = H_t + \alpha h$, you need to **re-compute** all the **weights** and then **re-normalize**. It is however straight-forward to show that the unnormalized weight \hat{w}_i is updated as $\hat{w}_i \leftarrow \hat{w}_i * e^{-\alpha h(\mathbf{x}_i)y_i}$ and that the normalizer Z becomes $Z \leftarrow Z * 2\sqrt{\epsilon(1-\epsilon)}$. Putting these two together we obtain the following multiplicative update rule: $w_i \leftarrow w_i \frac{e^{-\alpha h(\mathbf{x}_i)y_i}}{2\sqrt{\epsilon(1-\epsilon)}}$.

A few remarks:

- As long as H is negation closed (this means for every $h \in H$ we must also have $-h \in H$), it cannot be that the error $\epsilon > \frac{1}{2}$. The reason is simply that if h has error ϵ , it must be that $-h$ has error $1 - \epsilon$. So you could just **flip** h to $-h$ and obtain a classifier with **smaller** error. As h was found by minimizing the error, this is a contradiction.
- The inner loop can **terminate** as the **error $\epsilon = \frac{1}{2}$** , and in most cases it will converge to $\frac{1}{2}$ over time. In that case the latest weak learner h is only as good as a coin toss and cannot benefit the ensemble (therefore boosting terminates). Also note that if $\epsilon = \frac{1}{2}$ the step-size α would be zero.

Further analysis

Let us examine each one of these updates.

- The weight update:**

AdaBoost Pseudo-code

```

Input:  $\ell, \alpha, \{(\mathbf{x}_i, y_i)\}, \mathbb{A}$ 
 $H_0 = 0$ 
 $\forall i : w_i = \frac{1}{n}$ 
for  $t=0:T-1$  do
   $h = \operatorname{argmin}_h \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i$    [ $h = \mathbb{A}((w_1, \mathbf{x}_1, y_1), \dots, (w_n, \mathbf{x}_n, y_n))$ ]
   $\epsilon = \sum_{i:h(\mathbf{x}_i) \neq y_i} w_i$ 
  if  $\epsilon < \frac{1}{2}$  then
     $\alpha = \frac{1}{2} \ln \frac{1-\epsilon}{\epsilon}$ 
     $H_{t+1} = H_t + \alpha h$ 
     $\forall i : w_i \leftarrow \frac{w_i e^{-\alpha h(\mathbf{x}_i)y_i}}{2\sqrt{\epsilon(1-\epsilon)}}$ 
  else
    return ( $H_t$ )
  end
end
return ( $H_T$ )

```

Algorithm 1: AdaBoost in Pseudo-Code

$$\hat{w}_i \leftarrow \hat{w}_i * e^{-\alpha h(\mathbf{x}_i)y_i},$$

as, $h(\mathbf{x}_i)y_i$ is either $+1$ (if classified correctly by this weak learner) or -1 (otherwise), it follows that this weight update multiplies the weight w_i either by a factor $e^\alpha > 1$ if it was classified incorrectly (i.e. increases the weights), or by a factor $e^{-\alpha} < 1$ if it was classified correctly (i.e. decreases the weight).

- **Normalization update:**

$$Z \leftarrow Z * 2\sqrt{\epsilon(1-\epsilon)}.$$

Previously we established that the normalizer Z is identical to the loss. We can therefore use it to bound the loss function after T iterations:

$$\ell(H) = Z = n \prod_{t=1}^T 2\sqrt{\epsilon_t(1-\epsilon_t)},$$

(the factor n comes from the fact that the initial $Z_0 = n$, when all weights are $\frac{1}{n}$.) If we define $c = \max_t \epsilon_t$, we can establish

$$\ell(H) \leq n \left[2\sqrt{c(1-c)} \right]^T.$$

The function $c(1-c)$ is maximized at $c = \frac{1}{2}$. But we know that each $\epsilon_t < \frac{1}{2}$ (or else the algorithm would have terminated). Therefore $c(1-c) < \frac{1}{4}$ and we can re-write it as $c(1-c) = \frac{1}{4} - \gamma^2$, for some γ . This leaves us with

$$\ell(H) \leq n(1-4\gamma^2)^{\frac{T}{2}}.$$

In other words, the training loss is decreasing **exponentially**!

In fact, we can go even further and compute after how many iterations we must have zero training error. Note that the training loss is an *upper bound* on the training error (defined as $\sum_{i=1}^n \delta_{H(\mathbf{x}_i) \neq y_i}$) - simply because $\delta_{H(\mathbf{x}_i) \neq y_i} < e^{-y_i H(\mathbf{x}_i)}$ in all cases. We can then compute the number of steps required until the loss is less than 1, which would imply that not a single training input is misclassified.

$$n(1-4\gamma^2)^{\frac{T}{2}} < 1 \Rightarrow T > \frac{2 \log(n)}{\log(\frac{1}{1-4\gamma^2})}.$$

This is an amazing result. It shows that after $O(\log(n))$ iterations your training error must be zero. In practice it often makes sense to keep boosting even after you make no more mistakes on the training set.

Summary

Boosting is a great way to turn a *weak classifier* into a *strong classifier*. It defines a whole family of algorithms, including Gradient Boosting, AdaBoost, LogitBoost, and many others ... Gradient Boosted Regression Trees is one of the most popular algorithms for Learning to Rank, the branch of machine learning focused on learning ranking functions, for example for web search engines. A few additional things to know:

- The step size α is often referred to as *shrinkage*.
- Some people do not consider gradient boosting algorithms to be part of the boosting family, because they have no guarantee that the training error decreases exponentially. Often these algorithms are referred to as *stage-wise regression* instead.
- Inspired by Breiman's Bagging, *stochastic gradient boosting* subsamples the training data for each weak learner. This combines the benefits of bagging and boosting. One variant is to subsample only $n/2$ data points *without* replacement, which speeds up the training process.
- One advantage of boosted classifiers is that during test time the computation $H(\mathbf{x}) = \sum_{t=1}^T \alpha_t h_t(\mathbf{x}_t)$ can be stopped prematurely if it becomes clear which way the prediction goes. This is particularly interesting in search engines, where the exact ranking of results is typically only interesting for the top 10 search results. Stopping the evaluation of lower ranked search results can lead to tremendous speed ups. A similar approach is also used by the Viola-Jones algorithm to speed up face detection in images. Here, the algorithm scans regions of an image to detect possible faces. As almost all regions and natural images do not contain faces, there are huge savings if the evaluation can be stopped after just a few weak learners are evaluated. These classifiers are referred to as cascades, that spend very little time on the common case (no face), but more time on the rare interesting case (face). With this approach Viola and Jones were the first to solve face recognition in real-time on low performance hardware (e.g. cameras).
- AdaBoost is an extremely powerful algorithm, that turns any weak learner that can classify any weighted version of the training set with below 0.5 error into a strong learner whose training error decreases exponentially and that requires only $O(\log(n))$ steps until it is consistent.